

# **ECE/SIOC 228 Machine Learning for Physical Applications**

## **Lecture 12: Review and Decision Making in Physical Systems**

**Yuanyuan Shi**

Assistant Professor, ECE

University of California, San Diego



- Project Milestone Report is due next Monday, May 9<sup>th</sup>
- HW2 Written Part posted on Gradescope and Canvas.
- Due in 2 weeks on May 16<sup>th</sup>

# Course Outline

UC San Diego

## Part 0: Introduction & Fundamentals

- Lecture 2: Supervised Learning Setup
- Lecture 3: Neural Networks and Backpropagation
- Lecture 4: Optimization & Regularization in Deep Learning

## Part 1: Modeling and Learning in Physical Systems

- Lecture 5-6: Convolutional Neural Networks
- Lecture 7: Recurrent Neural Networks
- Lecture 8-9: Graph Neural Networks
- Lecture 10-11: Physics Informed Neural Networks and Neural Operator

# Lecture 2: Supervised Learning Setup

UC San Diego

- **Linear Regression**

Given  $\{x^{(i)}, y^{(i)}\}_{i=1}^N, x^{(i)} \in \mathbb{R}^n, y^{(i)} \in \mathbb{R}$ , we build a prediction model

$$\hat{y}^{(i)} = f(x^{(i)}) = b + \sum_{j=1}^n w_j x_j^{(i)} = w^T x^{(i)} + b$$

Cost function:

$$L(w, b) = \frac{1}{2} \sum_{i=1}^m (w^T x^{(i)} + b - y^{(i)})^2$$

- Gradient descent to find the best  $w$
- Why using squared loss? Maximum likelihood interpretation

- **Logistic Regression**

Given  $\{x^{(i)}, y^{(i)}\}_{i=1}^N, x^{(i)} \in \mathbb{R}^n, y^{(i)} \in \{0, 1\}$ , we build a prediction model

$$\hat{y}^{(i)} = \sigma(w^T x^{(i)} + b), \sigma(z) = \frac{1}{1 + e^{-z}}, \text{sigmoid function}$$

Cost function:

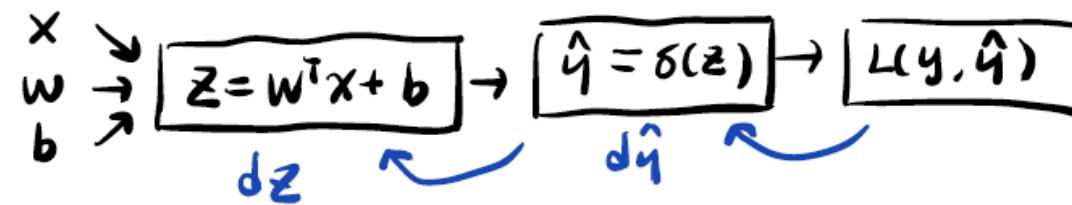
$$L(w, b) = - \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log (1 - \hat{y}^{(i)})]$$

# Lecture 3: Neural Networks and Backpropagation

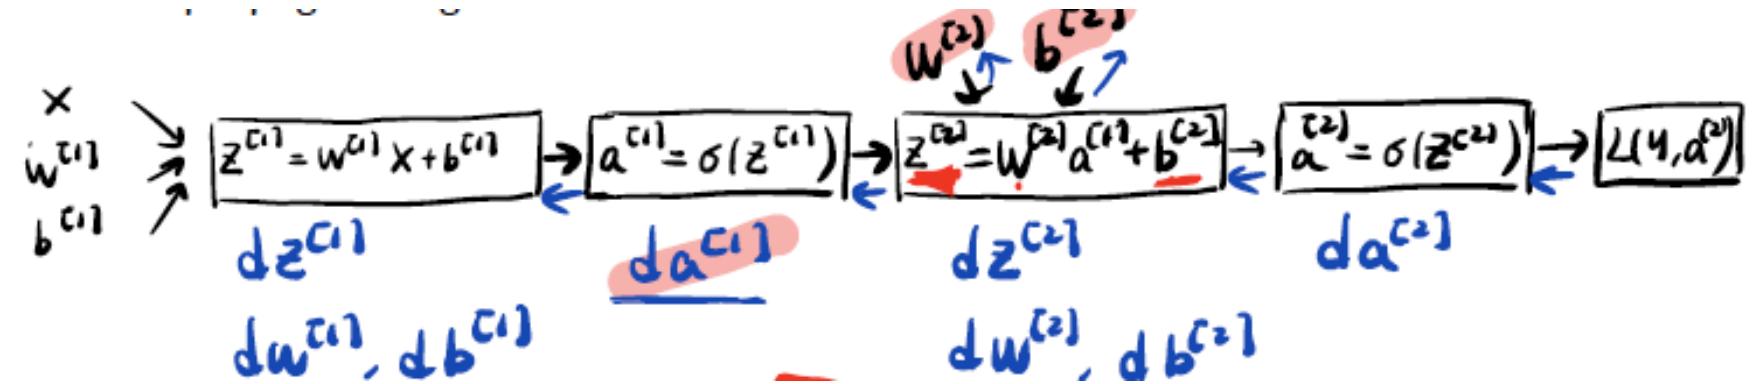
UC San Diego

- Computation Graph: Forward and Backward Computation

Logistic Regression



One hidden layer MLP

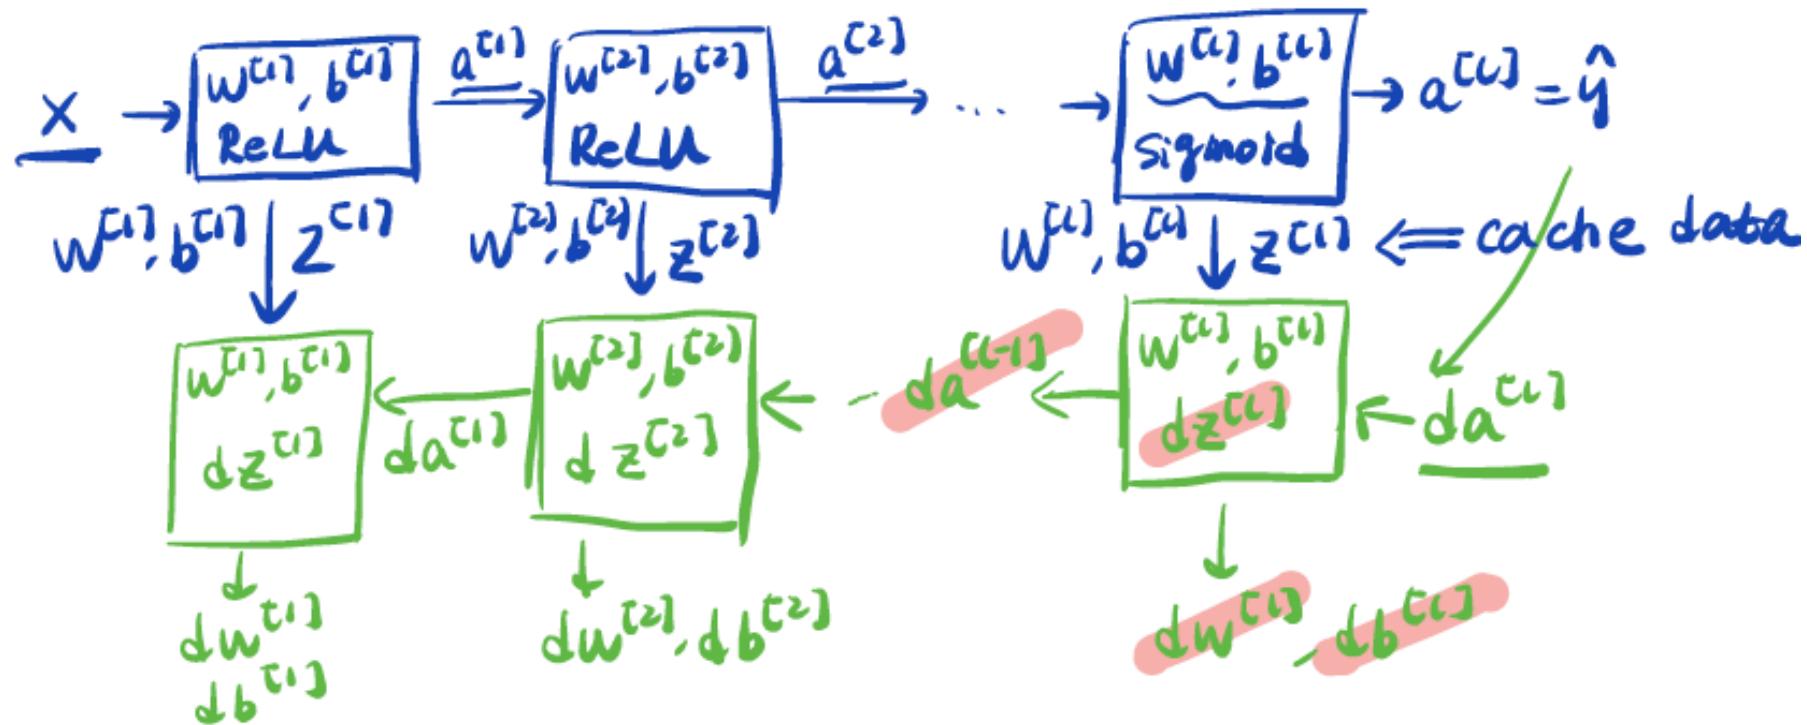


- Activation functions  $\sigma(\cdot)$ : sigmoid, tanh, ReLU, LeakyReLU; and their derivatives

# Lecture 3: Neural Networks and Backpropagation

UC San Diego

- Forward and Backward Computation of L-layer MLP



- Universal function approximation (1-layer NN) and why deep?

# Lecture 4: Optimization & Regularization in Deep Learning

UC San Diego

- Gradient descent / mini-batch gradient descent
- Gradient descent with momentum

$$V_{dw} = \beta V_{dw} + (1 - \beta) dw, V_{db} = \beta V_{db} + (1 - \beta) db$$

$$w^{(k+1)} = w^{(k)} - \eta V_{dw}, b^{(k+1)} = b^{(k)} - \eta V_{db}$$

- RMSprop (similar as Adagrad)

$$S_{dw} = \beta_2 S_{dw} + (1 - \beta_2)(dw)^2, S_{db} = \beta_2 S_{db} + (1 - \beta_2)(db)^2$$

$$w^{(k+1)} = w^{(k)} - \eta \frac{dw}{\sqrt{S_{dw}} + \varepsilon}, b^{(k+1)} = b^{(k)} - \eta \frac{db}{\sqrt{S_{db}} + \varepsilon}$$

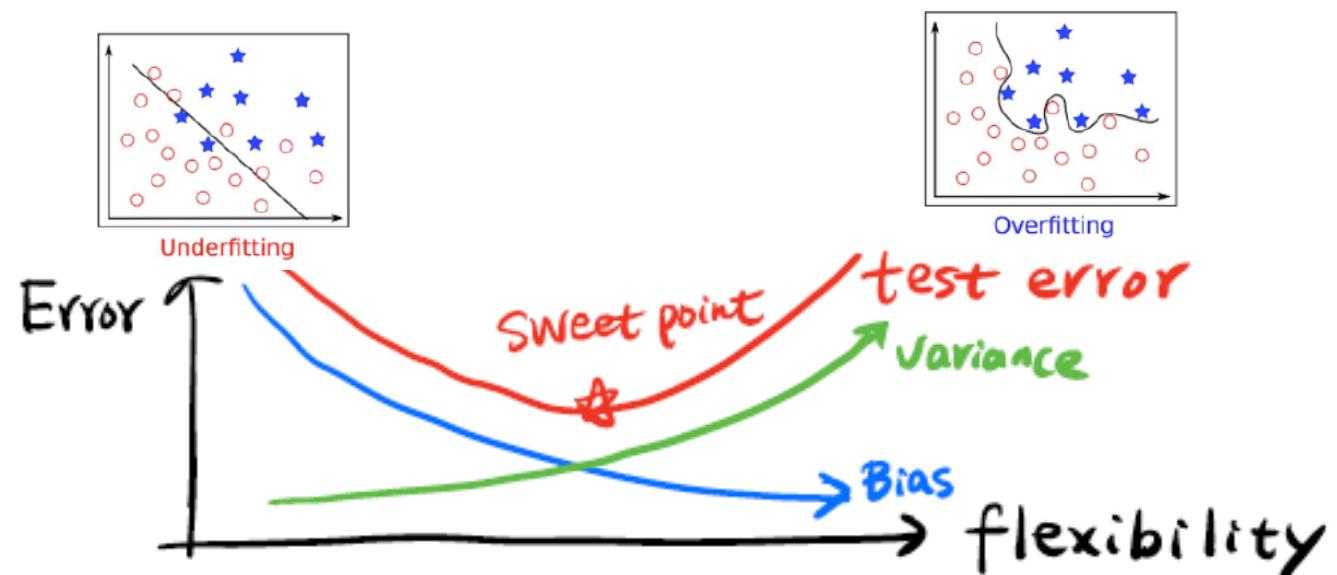
- Adam Optimizer

$$w^{(k+1)} = w^{(k)} - \eta \frac{V_{dw}}{\sqrt{S_{dw}} + \varepsilon}, b^{(k+1)} = b^{(k)} - \eta \frac{V_{db}}{\sqrt{S_{db}} + \varepsilon}$$

# Lecture 4: Optimization & Regularization in Deep Learning

UC San Diego

- Random weight initialization (and why we need it?)
- Learning rate decay
- Some recent optimization landscape conjecture in deep learning
  - No poor local minima
  - Need to escape saddle points effectively
- Regularization in Deep Learning
  - Bias and variance tradeoff
  - Dropout
  - **Batch normalization**
  - Early stopping
  - Weight regularization



# Course Outline

UC San Diego

## Part 0: Introduction & Fundamentals

- Lecture 2: Supervised Learning Setup
- Lecture 3: Neural Networks and Backpropagation
- Lecture 4: Optimization & Regularization in Deep Learning

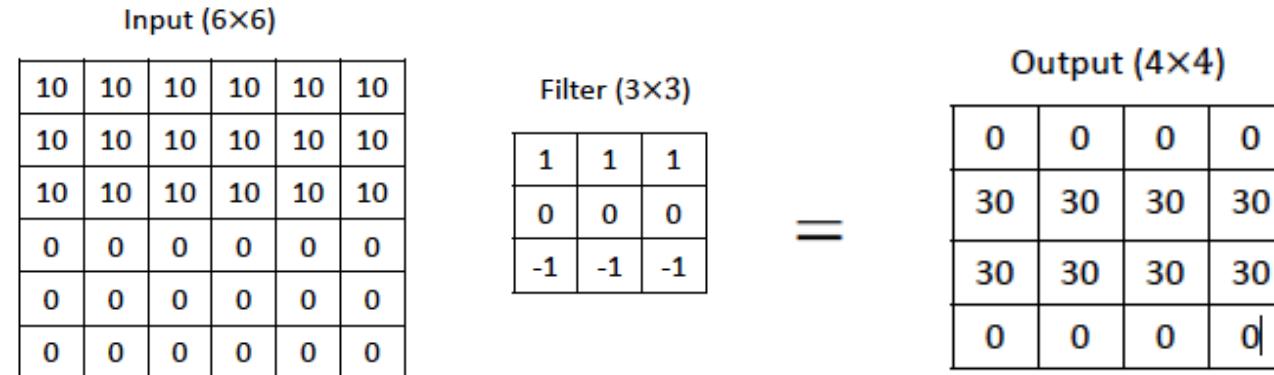
## Part 1: Modeling and Learning in Physical Systems

- Lecture 5-6: Convolutional Neural Networks
- Lecture 7: Recurrent Neural Networks
- Lecture 8-9: Graph Neural Networks
- Lecture 10-11: Physics Informed Neural Networks and Neural Operator

# Lecture 5-6: Convolutional Neural Networks

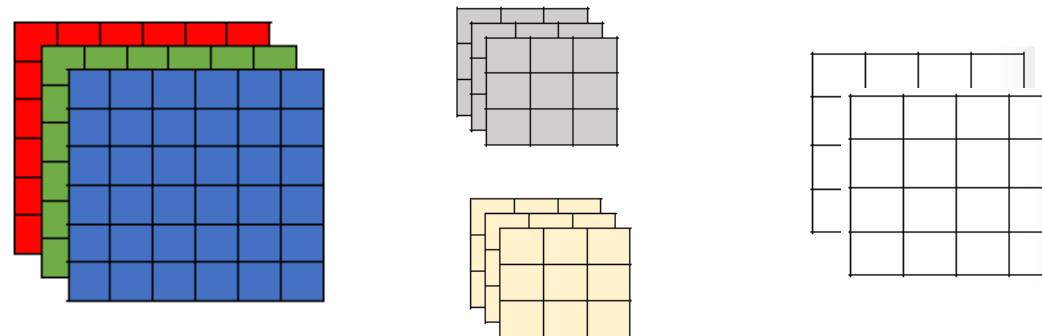
UC San Diego

- Convolution Layer
  - Filter ( $f$ )
  - Padding ( $p$ )
  - Stride ( $s$ )



$$\text{output size: } \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n + 2p - f}{s} + 1 \right\rfloor$$

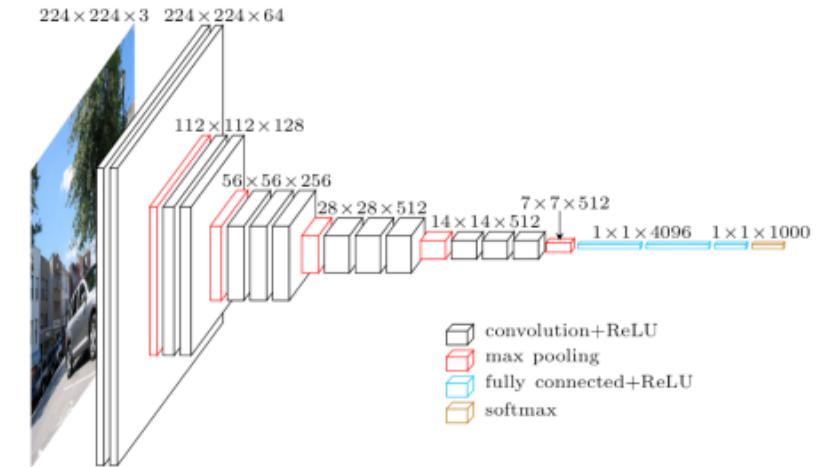
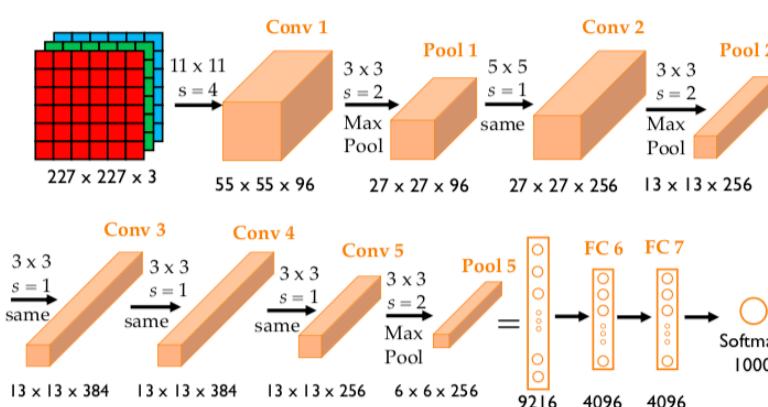
- Convolution over different channels
- More filters



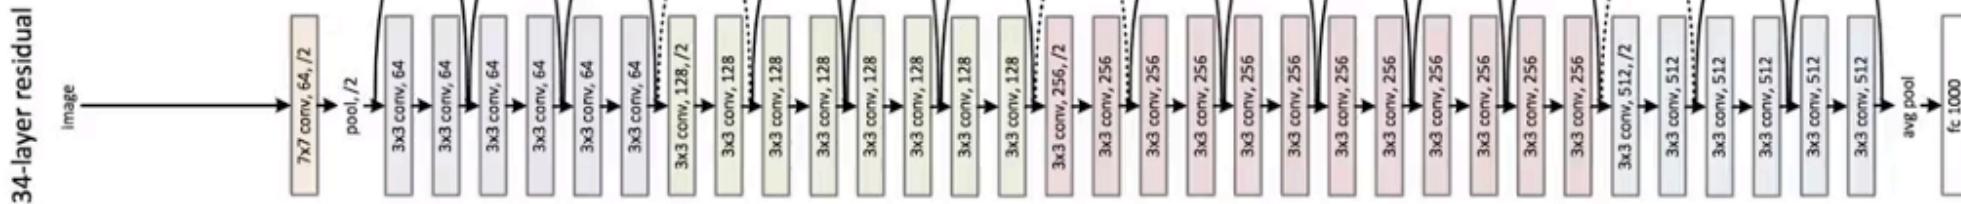
# Lecture 5-6: Convolutional Neural Networks

UC San Diego

- Pooling Layer
- Classic Networks
  - LeNet – 5
  - AlexNet
  - VGG-16
  - ResNet



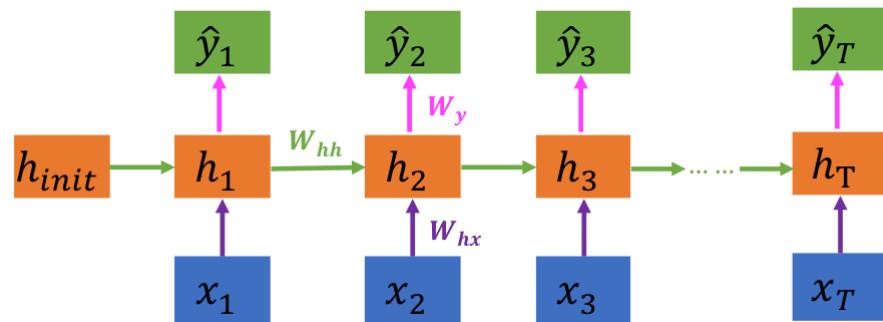
ResNet



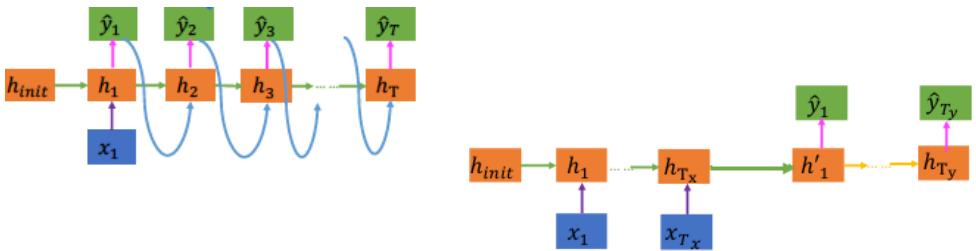
# Lecture 7: Recurrent Neural Networks

UC San Diego

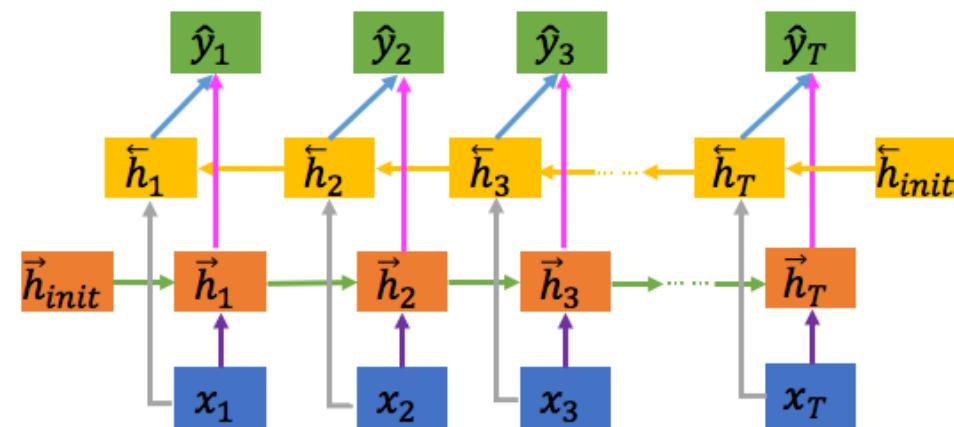
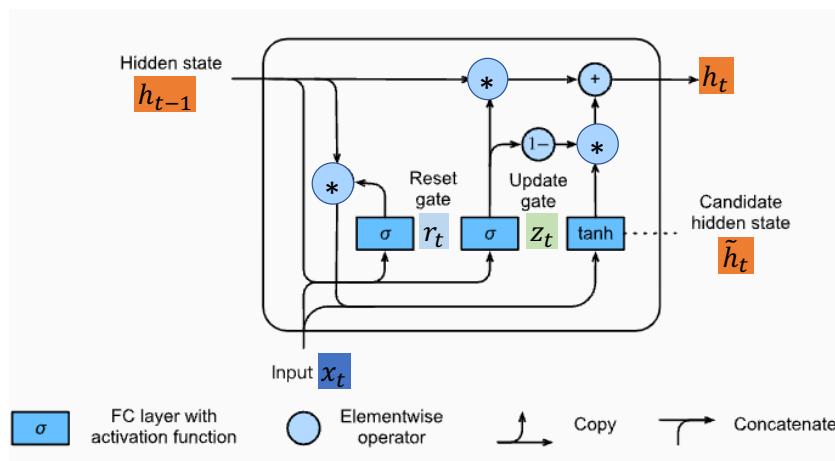
- Recurrent Neural Network: forward and backward computation through time



- One to many
- Many to one
- Many to many
- Many to many (input and output of different length)



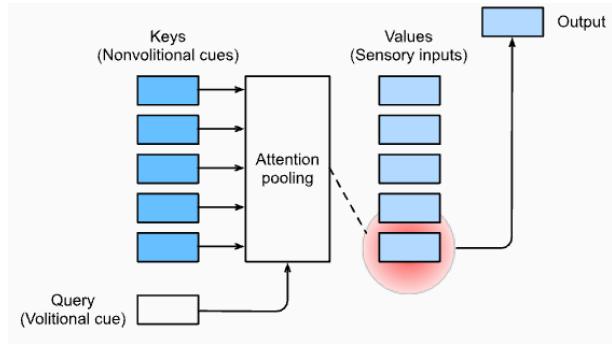
- Modern Architectures: Gated Recurrent Units (GRU), LSTM, Bidirectional RNN



# Lecture 7: Recurrent Neural Networks

UC San Diego

- Attention Mechanism



Query:  $q$ ; Keys and value pairs:  $(k_1, v_1), (k_2, v_2), \dots, (k_m, v_m)$

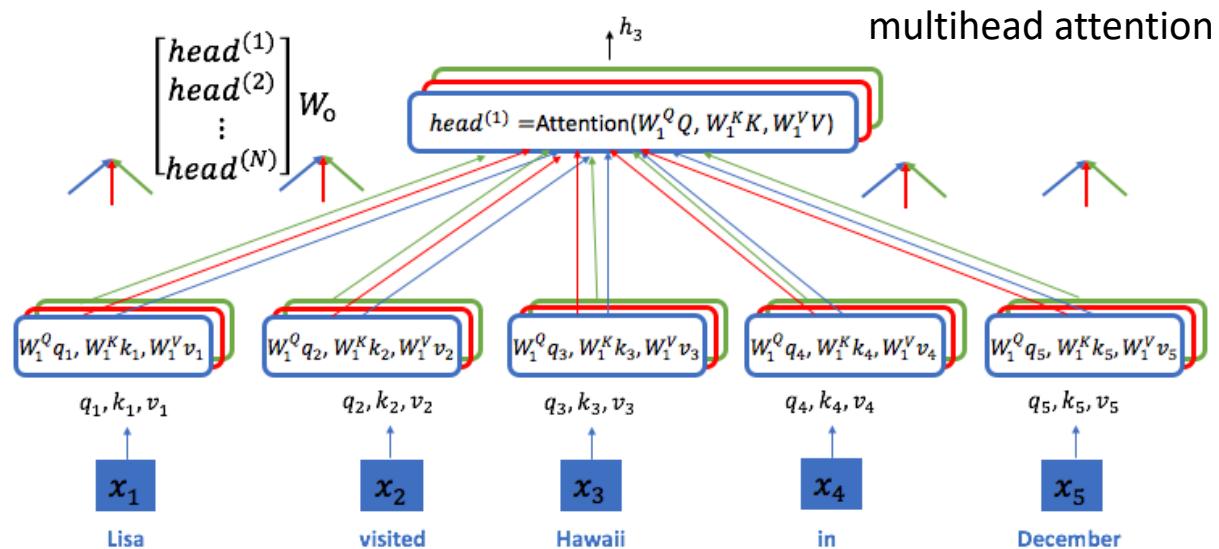
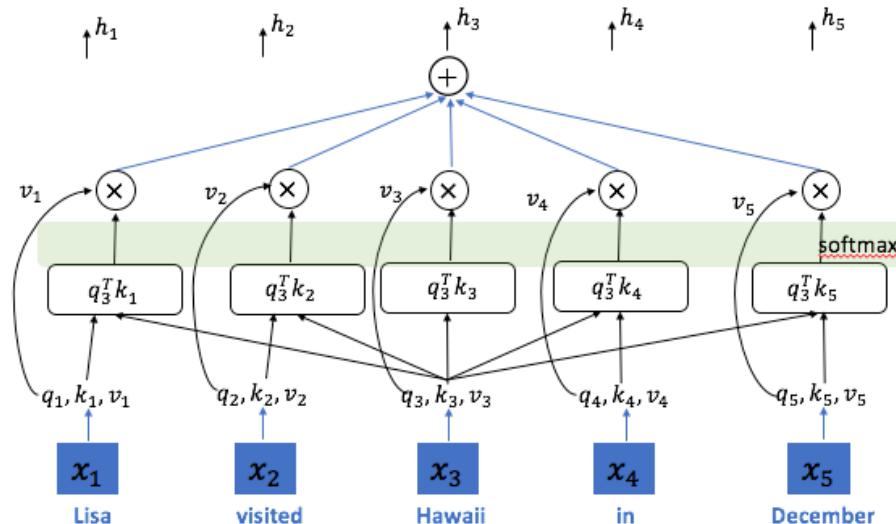
$$f(q) = \sum_{i=1}^m \alpha(q, k_i) v_i$$

$$\alpha(q, k_i) = \frac{\exp(a(q, k_i))}{\sum_{j=1}^m \exp(a(q, k_j))} \in (0, 1), \text{ } a(\cdot, \cdot) \text{ is the attention score function.}$$

$$a(q, k_i) = \frac{k_i^T q}{\sqrt{d_k}} \text{ or } a(q, k_i) = W_a \sigma(W_q q + W_k k_i)$$

- Transformer

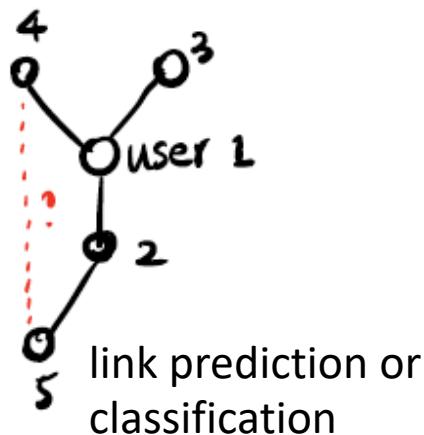
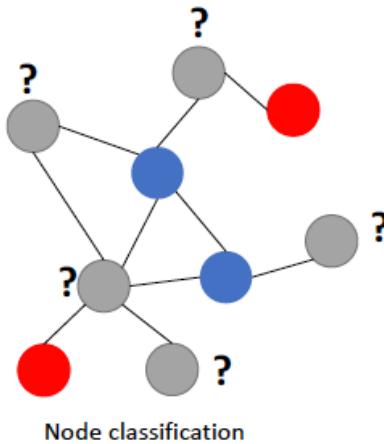
self attention



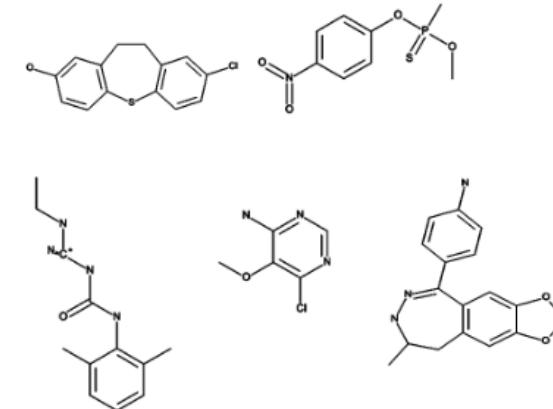
# Lecture 8-9: Graph Neural Networks

UCSanDiego

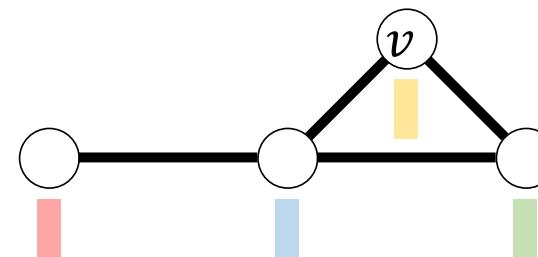
- Node embedding is the key for different graph tasks



Graph-level task: toxic molecule or not?



- 3 popular graph neural network architectures
  - Graph convolutional networks
  - Graph Sample and Aggregate (GraphSAGE)
  - Graph Attention Networks



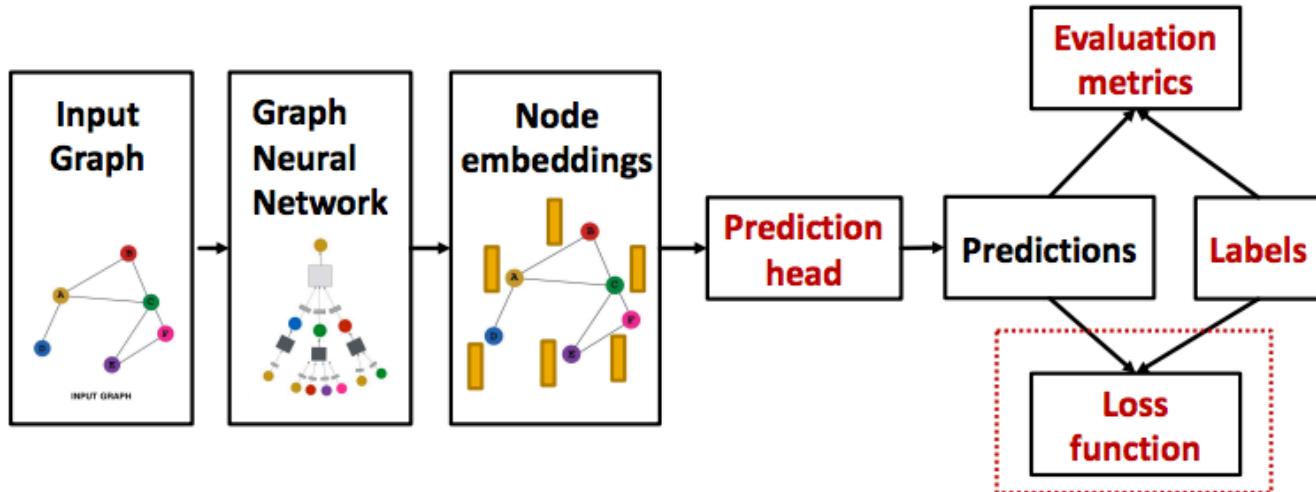
$$m_v = \text{Agg}(\{\text{blue bar}, \text{green bar}\})$$
$$h_v^{(k)} = f^{(k)}[\text{combine}(x_v, m_v)]$$

1-layer of GNN

# Lecture 8-9: Graph Neural Networks

UC San Diego

- Pipeline: end-to-end training of GNN and prediction models



Some example applications:

Node classification

$$\hat{y}_v = \text{prediction model}(h_v^{(k)})$$

Loss function:

$$L(\hat{y}_v, y_v) = - \sum_{c \in \text{all classes}} y_v^{(c)} \log \hat{y}_v^{(c)}$$

Edge prediction

$$\hat{y}_{uv} = \text{prediction model}(h_v^{(k)}), P_{uv} = \text{sigmoid}(\hat{y}_v^T \hat{y}_u)$$

Loss function:

$$L(\hat{y}_v, \hat{y}_u, e_{uv}) = -e_{uv} \log P_{uv} - (1 - e_{uv}) \log (1 - P_{uv})$$

Graph level task

$$G^{(k)} = \text{Agg/Pool}[h_{v1}^{(k)}, h_{v2}^{(k)}, \dots, h_{vN}^{(k)}]$$

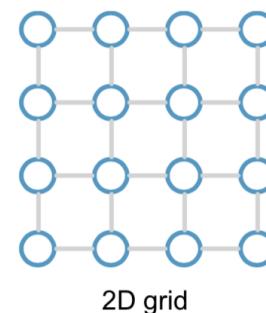
$$\hat{y}_G = \text{prediction model}(G^{(k)})$$

$$\text{Loss function: } L(\hat{y}_G, y_G) = - \sum_{c \in \text{all classes}} y_G^{(c)} \log \hat{y}_G^{(c)}$$

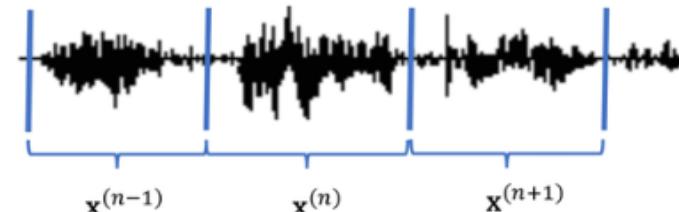
# CNN, RNN, GNN

UC San Diego

basic convolutional neural networks, recurrent networks and graph neural networks can all be viewed as special cases of standard neural networks (MLP) with shared weights and skip connections

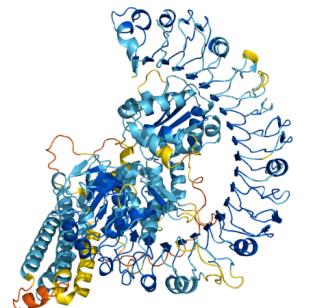
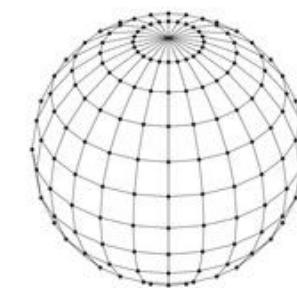


CNN operates on 2D spatial structure



RNN operates on 1D temporal structure

- GNN can operate on spatial data with complex geometry structure

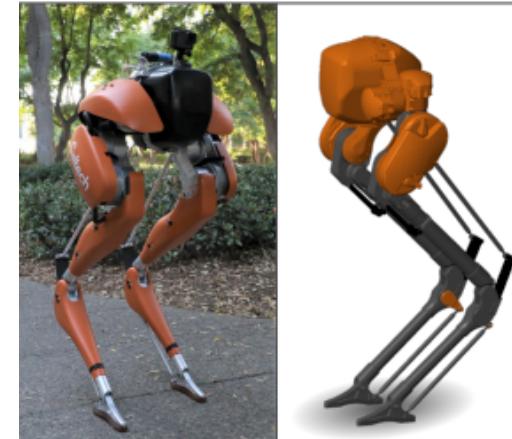


spherical

# Lecture 10-11: Physics Informed Neural Networks and Neural Operator

UC San Diego

- Many physical systems are governed or described by ODEs and PDEs



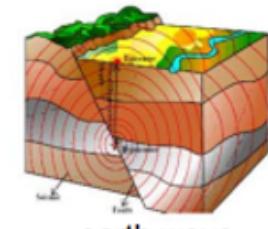
$$1D \text{ Wave Equation: } \frac{\partial^2 u}{\partial t^2} = c^2 \frac{\partial^2 u}{\partial x^2}$$



water wave



air wave



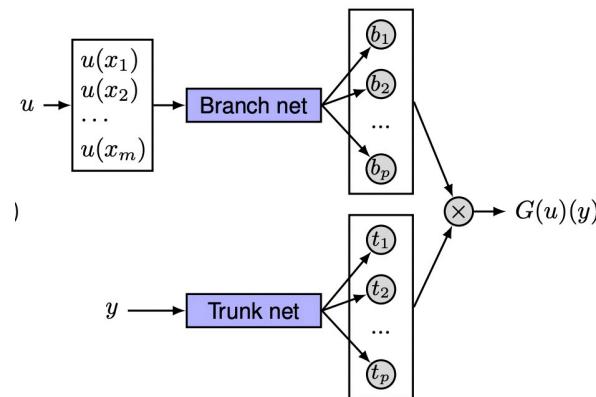
earth wave

- If we know the form of the differential equations
  - Could use PINN to solve and predict how the system trajectory will evolve
  - Or use PINN for learning/identifying the unknown parameters in the systems

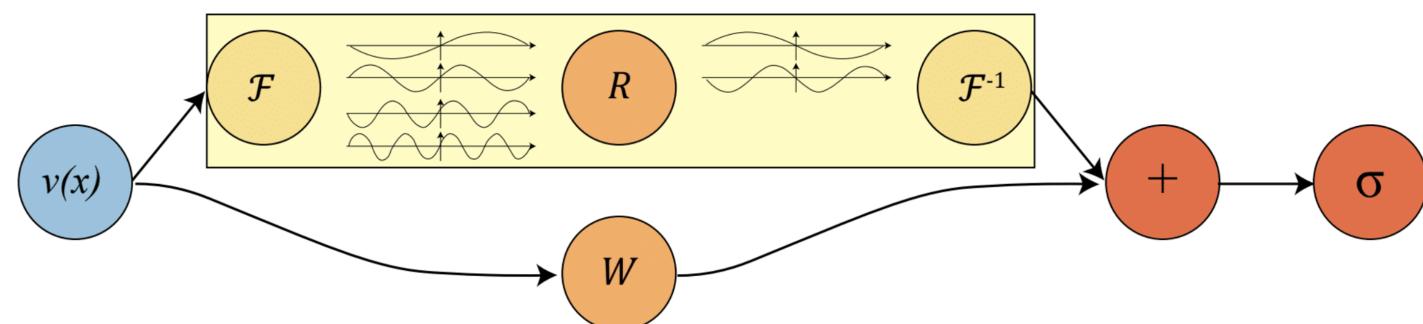
# Lecture 10-11: Physics Informed Neural Networks and Neural Operator

UC San Diego

- If we only have data collected from the physical systems (governed or described by ODEs and PDEs), we can learn a neural operator to model and predict the system behaviors.



DeepONet



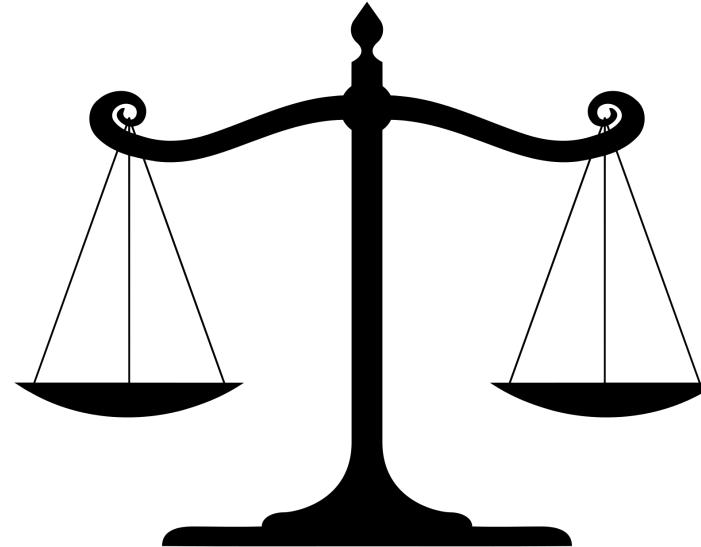
Fourier Neural Operator

# Summary: Modeling and Learning in Physical Systems

UC San Diego

## Priors: things assumed beforehand

- Which model to use?
- Which prior knowledge (physics law/constraints/...) to enforce?



## Learning: things extracted from data

- Basic supervised learning recipe
- Gather training data; decide model class
  - Define loss function
  - Improve the model by minimizing loss (optimization problem – often use gradient descent in deep learning)

- Balance between prior and learning:
  - Strong prior, minimal learning: can lead to bias if the prior isn't aligned with the data
  - Weak prior, much learning: may lead to high variance if overfitting to the noise in the data
- Choose the priors and collect data to obtain a model that achieves desired performance on your task

# Course Outline

UC San Diego

## Part 0: Introduction & Fundamentals

- Lecture 2: Supervised Learning Setup
- Lecture 3: Neural Networks and Backpropagation
- Lecture 4: Optimization & Regularization in Deep Learning

## Part 1: Modeling and Learning in Physical Systems

- Lecture 5-6: Convolutional Neural Networks
- Lecture 7: Recurrent Neural Networks
- Lecture 8-9: Graph Neural Networks
- Lecture 10-11: Physics Informed Neural Networks and Neural Operator

## Part 2: Decision Making in Physical Systems

# Decision Making in Physical Systems

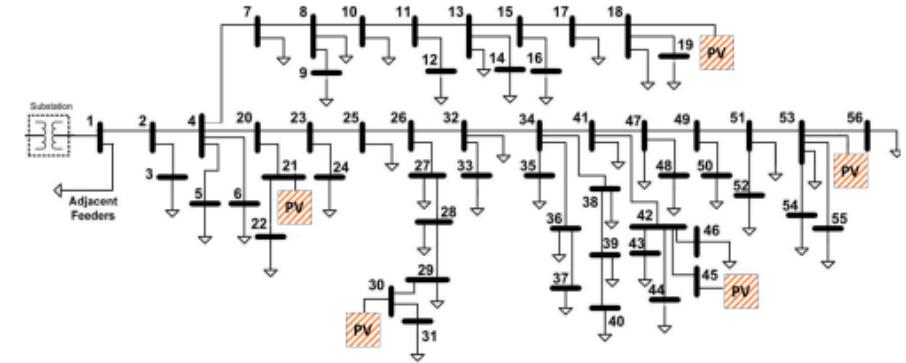
UC San Diego



Robots: decide where to move the arm



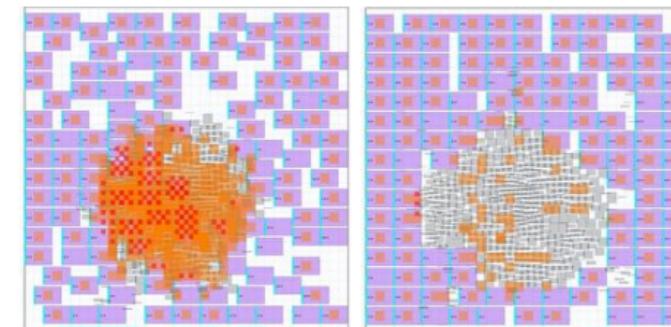
Autonomous vehicles: to accelerate or slow down?



Power Grid: decide how much power to generate



Supply chain: decide how much inventory to replenish



Chip Placement

# Uncertainty in the Physical Systems

UC San Diego



**Robots:** decide where the arm move to, but actuators can fail, hit unseen obstacles, etc



**Renewable generators:** decide how much to produce, but sun does not always shine (cloud) and wind does not always blow



**Inventory management:** decide what to replenish, but don't know the customer demand for various products

# Outline

UC San Diego

- **Markov Decision Process**
- Value Iteration
- Policy Evaluation

# Overview of Concepts

UC San Diego

- Markov Decision Process
  - State
  - Action
  - Transition Probability
  - Reward (and discount factor)
- Policy
- Return
- Value function

# Markov Property

UC San Diego

- **Markov property:** A state  $S_t$  is Markov if and only if

$$P[S_{t+1}|S_t] = P[S_{t+1}|S_1, \dots, S_t]$$

- The future is independent of the past given the present
- **State transition probability:** from a Markov state  $s$  to next state  $s'$ , the state transition probability is defined by:

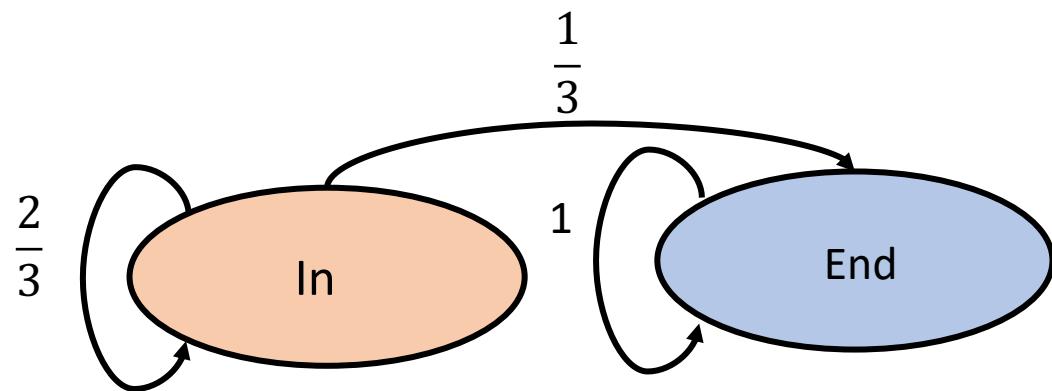
$$P_{ss'} = P[S_{t+1} = s' | S_t = s]$$

- **State transition matrix**

# Markov Process & Example

UC San Diego

- A Markov Process (or Markov Chain) in a tuple  $\langle S, P \rangle$ 
  - $S$  is a finite set of states
  - $P$  is a state transition probability matrix,  $P_{ss'} = P[S_{t+1} = s' | S_t = s]$

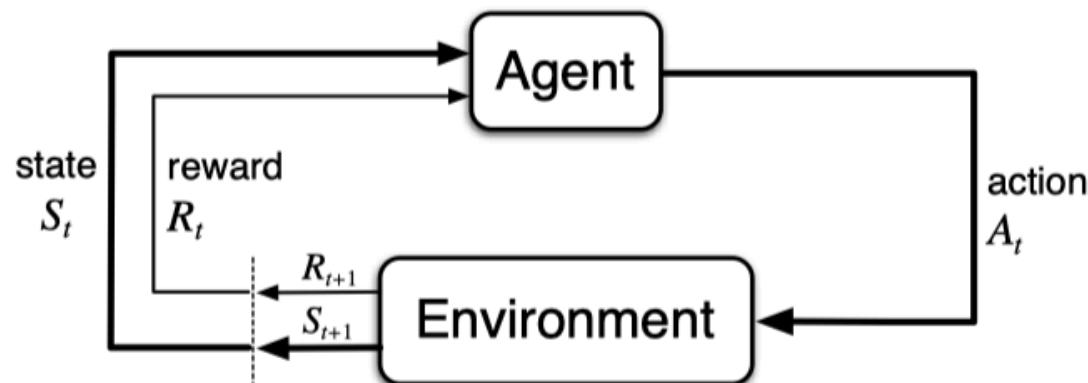


Dice Game Markov Chain

- For each round  $r = 1, 2, \dots$
- Roll a 6-sided dice
  - If the dice results in 1 or 2, game ends
  - Otherwise, game continues to the next round

# Markov Decision Process

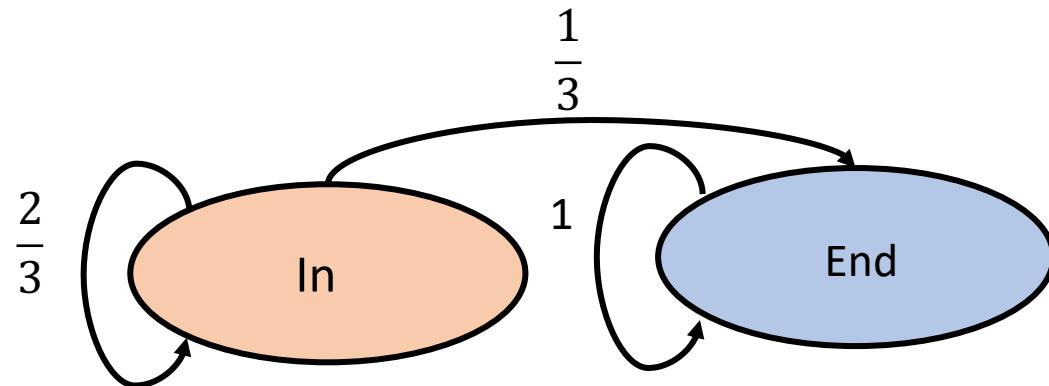
- A Markov Decision Process in a tuple  $\langle S, A, P, R, \gamma \rangle$ 
  - $S$  is a finite set of states
  - $A$  is a finite set of actions
  - $P$  is a state transition probability matrix,  $P_{ss'}^a = P[S_{t+1} = s' | S_t = s, A_t = a]$
  - $R$  is a reward function,  $R_s^a = R(R_{t+1} | S_t = s, A_t = a)$
  - $\gamma$  is a discount factor  $\gamma \in [0, 1]$



**Figure 3.1:** The agent–environment interaction in a Markov decision process.

# Markov Decision Process Example

UC San Diego



Dice Game Markov Chain

For each round  $r = 1, 2, \dots$

- You choose **stay** or **quit**
- If quit, you get \$10 and the game ends
- If stay, you get \$4 and then we roll a 6-sided dice
  - If the dice results in 1 or 2, game ends
  - Otherwise, game continues to the next round

# Two Sets of Notations?

UC San Diego

$s_t$  – state

$a_t$  – action

$r(s, a)$  – reward function

$x_t$  – state

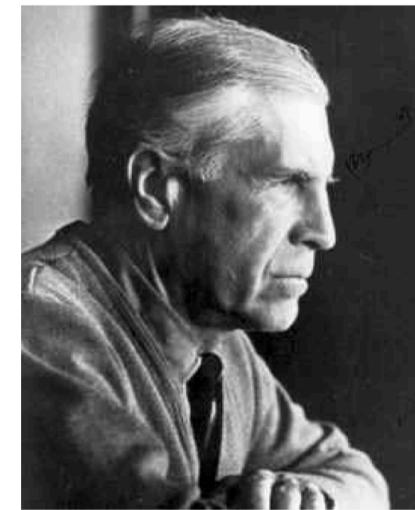
$u_t$  – action

$c(x, u)$  – cost function

$$r(s, a) = -c(x, u)$$

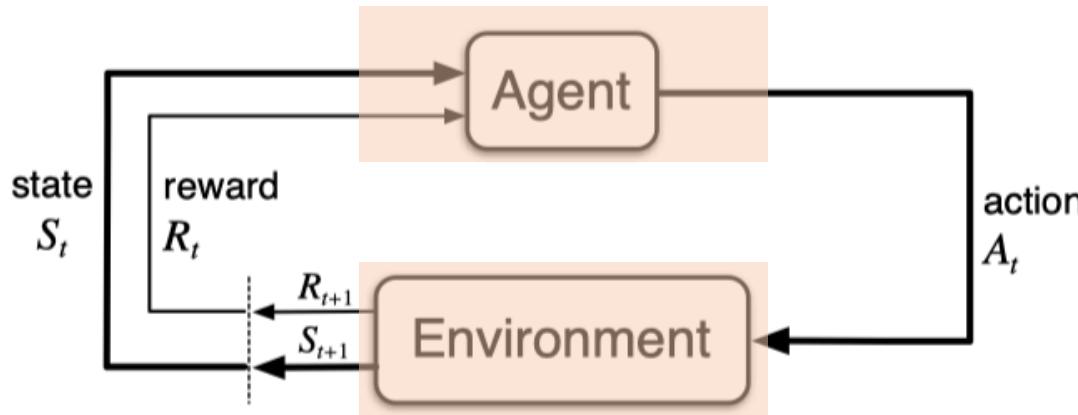


Richard Bellman



Lev Pontryagin

# Environment Model



**Figure 3.1:** The agent–environment interaction in a Markov decision process.

- A **model** predicts what the environment will do next
- $\mathcal{P}$  predicts the next state
- $\mathcal{R}$  predicts the next (immediate) reward, e.g.

$$\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$$

$$\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$$

- A policy  $\pi$  is a mapping from each state  $s \in S$  to an action  $a \in A$

$$\pi[a|s] = P[A_t = a|S_t = s]$$

- A policy fully defines the behavior of an agent
- In MDP, policies only depend on the current state  $S_t$ , not the history
- Policy is stationary (time independent),  $A_t \sim \pi(\cdot | S_t), \forall t > 0$

## Example: policy "stay"

For each round  $r = 1, 2, \dots$

- You choose stay or quit
- If quit, you get \$10 and the game ends
- If stay, you get \$4 and then we roll a 6-sided dice
  - If the dice results in 1 or 2, game ends
  - Otherwise, game continues to the next round

Sample episodes starting from  $S_1 = \text{In}$

- [In (stay), End] \$4
- [In (stay), In (stay), In (stay), End] \$12
- [In (stay), In (stay), End] \$8
- [In (stay), In (stay), In (stay), In (stay), End] \$16
- .....

- A policy  $\pi$  is a mapping from each state  $s \in S$  to an action  $a \in A$

$$\pi[a|s] = P[A_t = a|S_t = s]$$

- A policy fully defines the behavior of an agent
- In MDP, policies only depend on the current state  $S_t$ , not the history
- Policy is stationary (time independent),  $A_t \sim \pi(\cdot | S_t), \forall t > 0$

## Example: policy "quit"

For each round  $r = 1, 2, \dots$

- You choose stay or quit
- If quit, you get \$10 and the game ends
- If stay, you get \$4 and then we roll a 6-sided dice
  - If the dice results in 1 or 2, game ends
  - Otherwise, game continues to the next round

Sample episodes starting from  $S_1 = \text{In}$

- [In (quit), End] \$10

# Return

UC San Diego

- The return  $G_t$  is the total discounted reward from time step  $t$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- The discount  $\gamma \in [0, 1]$  is the present value of future rewards
- The value of receiving reward  $R$  after  $k + 1$  time-steps is  $\gamma^k R$
- This values immediate reward above delayed reward
  - $\gamma$  close to 0 leads to "greedy" evaluation (only care about current step)
  - $\gamma$  close to 1 leads to "far-sighted" evaluation

# Return

UCSanDiego

- The return  $G_t$  is the total discounted reward from time step  $t$

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

For each round  $r = 1, 2, \dots$

- You choose stay or quit
- If quit, you get \$10 and the game ends
- If stay, you get \$4 and then we roll a 6-sided dice
  - If the dice results in 1 or 2, game ends
  - Otherwise, game continues to the next round

Discount  $\gamma = 1$ : [In (stay), In (stay), In (stay), In (stay), End]  $4 + 4 + 4 + 4 = \$16$

Discount  $\gamma = 0$ : [In (stay), In (stay), In (stay), In (stay), End]  $4 + 0 \cdot 4 + 0 \cdot 4 + 0 \cdot 4 = \$4$

Discount  $\gamma = 1/2$ : [In (stay), In (stay), In (stay), In (stay), End]  $4 + 0.5 \cdot 4 + 0.5^2 \cdot 4 + 0.5^3 \cdot 4 = \$7.5$

# Why Discount?

- Mathematically convenient to use discount rewards
- Animal/human behavior shows preference for immediate reward
- If the reward is financial, immediate rewards may earn more interest than delayed rewards
- Sometimes, it is possible to use undiscounted reward if all possible episodes are finite
- .....

# Value function

UC San Diego

The **state-value function  $v_\pi(s)$  of a policy  $\pi$**  in an MDP is the expected return starting from state  $s$ , and then following policy  $\pi$

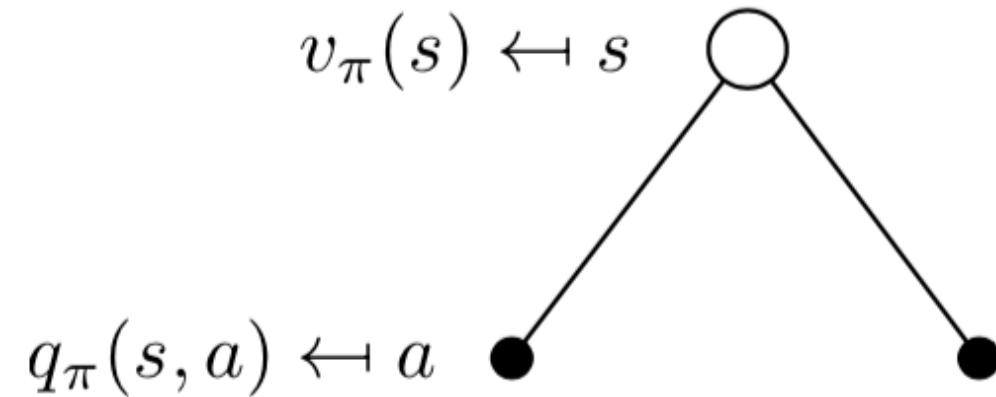
$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$$

The **state-action value function  $q_\pi(s, a)$  of a policy  $\pi$**  (sometimes also refer to as action value function) is the expected return starting from state  $s$ , taking action  $a$ , and then following policy  $\pi$

$$q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$

# Value function

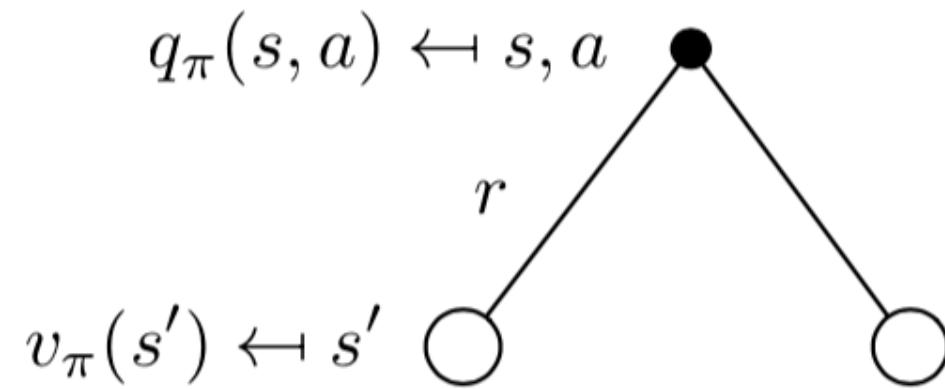
UC San Diego



$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) q_\pi(s, a)$$

# Value function

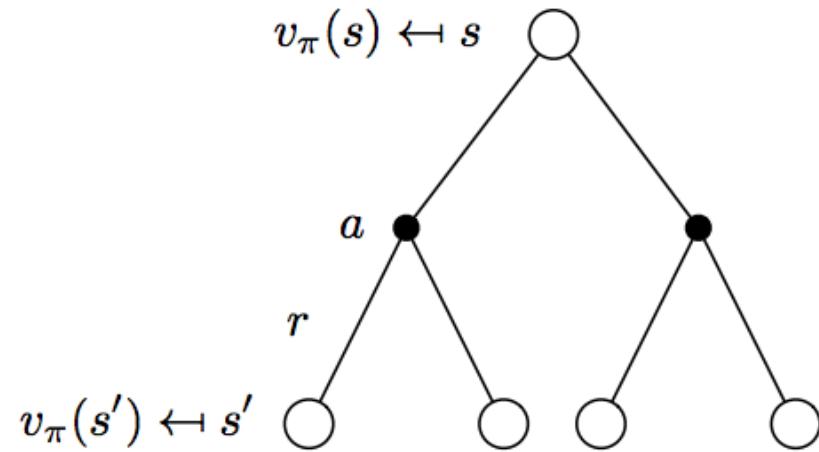
UC San Diego



$$q_{\pi}(s, a) = \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_{\pi}(s')$$

# Value function

UC San Diego



$$v_\pi(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_\pi(s') \right)$$

# Summary

- Markov Decision Process  $\langle S, A, P, R, \gamma \rangle$ 
  - State
  - Action
  - Transition Probability
  - Reward (and discount factor)
- Policy
- Return
- Value function: state value function  $v_\pi(s)$ , action value function  $q_\pi(s, a)$